

Unit 12


Text Embeddings

EE-UY 4563/EL-GY 9143: INTRODUCTION TO MACHINE LEARNING
PROF. SUNDEEP RANGAN

Learning Objectives

- ❑ Describe the function of a **text embedding**
- ❑ Describe the role of **tokenization** in a text embedding
- ❑ Perform **Byte-Pair Encoding** on simple corpusS
- ❑ Describe **attention architecture** in modern LLMs
- ❑ Encode text with a **pre-trained text embedding**
- ❑ Use the text embedding for simple **downstream tasks** such as classification

Outline

- 
- Text Embeddings and Tokenization
 - Modern LLMs at a Glance
 - Using a pre-trained LLM

What is a Text Embedding?


 **Sierra P**
34 reviews · 18 photos

★★★★★ a month ago

Dine in | Dinner | \$50–100

Wonderful establishment. Came here for a friend's birthday and the treatment was amazing. They even surprised her with a birthday cake which was so unexpected. I absolutely adore the waiter and staff for their kindness and will be returning. The place is vegan friendly and food is delicious with a decent price. The atmosphere is extremely elegant and great to celebrate her 21st birthday in a toned manner. Everything 10/10 and we came at an amazing hour where the restaurant was empty on a weekday. Loveee.

I ordered : mash potatoes, asparagus, and a cocktail. (Vegan)



- ❑ ML models work on **numeric data**
- ❑ Input is typically a **vector**
- ❑ **Embedding:** Maps text to a vector
- ❑ Used by **downstream tasks**
 - E.g., classification

Three Key Goals

❑ Compactness

- Low dimensionality relative to the input space
- Enables efficient storage and faster downstream computation

❑ Expressiveness


- Capture the salient features of the input
- Preserve semantic, syntactic, and contextual information

❑ Semantic Similarity

- Texts with similar meaning should have nearby vectors
- Enables clustering, retrieval, and generalization

Sierra P
34 reviews · 18 photos
★★★★★ a month ago
Dine in | Dinner | \$50–100
Wonderful establishment. Came here for a friend's birthday and the treatment was amazing. They even surprised her with a birthday cake which was so unexpected. I absolutely adore the waiter and staff for their kindness and will be returning. The place is vegan friendly and food is delicious with a decent price. The atmosphere is extremely elegant and great to celebrate her 21st birthday in a toned manner. Everything 10/10 and we came at an amazing hour where the restaurant was empty on a weekday. Loveee.
I ordered : mash potatoes, asparagus, and a cocktail. (Vegan)

Embedding



Vector

Typical Embedding Flow

Original text


 **Sierra P**
34 reviews · 18 photos
★★★★★ a month ago
Dine in | Dinner | \$50–100
Wonderful establishment. Came here for a friend's birthday and it was amazing. They even surprised her with a birthday cake which I unexpected. I absolutely adore the waiter and staff for their kindr be returning. The place is vegan friendly and food is delicious wit price. The atmosphere is extremely elegant and great to celebrate birthday in a toned manner. Everything 10/10 and we came at an hour where the restaurant was empty on a weekday. Loveee.

I ordered : mash potatoes, asparagus, and a cocktail. (Vegan)


Tokenization


Token-level
embedding

'Wonderful',
'establishment'.
'Came'
'here',
...
...



Document
embedding



Character and Word Tokenization


- Early systems used one of two simple possibilities for tokens: Character or Word
- **Character level:** Each character is a token
- **Word level:** Each word is a token
- Consider example: “The food was delicious”

Encoding type	Token	Robustness	Keeps semantic meaning
Character-level	‘T’, ‘h’, ‘e’, ‘ ’, ‘f’, ...	Simple, works on all text	No semantic meaning
Word-level	‘The’, ‘food’, ‘was’, ‘delicious’	Only if all words are in “vocabulary”	Yes, with proper training

Word Embeddings

- ❑ Used with word-level tokenization
- ❑ Assign each word token to a vector
- ❑ Training goal:
 - Assign similar words similar embeddings
 - Trained by looking a “context word”
 - Context word = word close to target

- ❑ Example: GloVe
 - Global Vectors for Word Representation
 - Stanford 2014



Limitations of Word Embeddings

- ❑ Word embeddings were widely-used for many years (up to about 2016)
- ❑ But hit many limitations
- ❑ Vocabulary size must grow
- ❑ Even with large vocabularies may out-of-vocabulary words
- ❑ Also, words lose context in sentence

<https://nlp.stanford.edu/projects/glove/>

Download pre-trained word vectors

- Pre-trained word vectors. This data is made available under the [Public Domain Dedication and License v1.0](http://www.opendatacommons.org/licenses/pddl/1.0/) whose full text can be found at: <http://www.opendatacommons.org/licenses/pddl/1.0/>.
 - **NEW!!** 2024 Dolma (220B tokens, 1.2M vocab, uncased, 300d vectors, 1.6 GB download): [glove.2024.dolma.300d.zip](https://nlp.stanford.edu/projects/glove/glove.2024.dolma.300d.zip)
 - **NEW!!** 2024 Wikipedia + Gigaword 5 (11.9B tokens, 1.2M vocab, uncased, 300d vectors, 1.6 GB download): [glove.2024.wikigiga.300d.zip](https://nlp.stanford.edu/projects/glove/glove.2024.wikigiga.300d.zip)
 - **NEW!!** 2024 Wikipedia + Gigaword 5 (11.9B tokens, 1.2M vocab, uncased, 200d vectors, 1.1 GB download): [glove.2024.wikigiga.200d.zip](https://nlp.stanford.edu/projects/glove/glove.2024.wikigiga.200d.zip)
 - **NEW!!** 2024 Wikipedia + Gigaword 5 (11.9B tokens, 1.2M vocab, uncased, 100d vectors, 560 MB download): [glove.2024.wikigiga.100d.zip](https://nlp.stanford.edu/projects/glove/glove.2024.wikigiga.100d.zip)
 - **NEW!!** 2024 Wikipedia + Gigaword 5 (11.9B tokens, 1.2M vocab, uncased, 50d vectors, 290 MB download): [glove.2024.wikigiga.50d.zip](https://nlp.stanford.edu/projects/glove/glove.2024.wikigiga.50d.zip)
 - [Wikipedia 2014 + Gigaword 5](https://nlp.stanford.edu/projects/glove/glove.6B.zip) (6B tokens, 400K vocab, uncased, 50d, 100d, 200d, & 300d vectors, 822 MB download): [glove.6B.zip](https://nlp.stanford.edu/projects/glove/glove.6B.zip)
 - Common Crawl (42B tokens, 1.9M vocab, uncased, 300d vectors, 1.75 GB download): [glove.42B.300d.zip](https://nlp.stanford.edu/projects/glove/glove.42B.300d.zip)
 - Common Crawl (840B tokens, 2.2M vocab, cased, 300d vectors, 2.03 GB download): [glove.840B.300d.zip](https://nlp.stanford.edu/projects/glove/glove.840B.300d.zip)
 - Twitter (2B tweets, 27B tokens, 1.2M vocab, uncased, 25d, 50d, 100d, & 200d vectors, 1.42 GB download): [glove.twitter.27B.zip](https://nlp.stanford.edu/projects/glove/glove.twitter.27B.zip)

• Download for downstream Twitter data

Sub-Word Tokenization

- ❑ Sub-word tokenization: Overcomes limitations of character and word-level tokenization
- ❑ Became widely-used around 2016-2018, especially with BERT
- ❑ Consider example: "**The food was unicornishly delicious :)**"

Encoding type	Token	Procs and cons
Character-level	'T', 'h', 'e', ' ', 'f', ...	Simple, works on all text, but no semantic meaning
Word-level	'The', 'food', 'was', OOV, 'delicious', OOV	Keeps semantic meaning of words, but may be out of vocabulary items
Sub-word	'The', 'food', 'was', 'unicorn', 'ish', 'ly', 'delicious'	Applies to all words and keeps semantic meaning

ASCII Characters

□ American Standard Code for Information Interchange

□ 1 byte (8 bits) for each character

- Encodes 128 characters
- 95 English characters (upper and lower case)
- 10 digits (0 to 9)
- Punctuation characters
- 33 control characters (CR, LF, Tab, ...)

□ Limitations

- Non-English characters
- Accented letters (e.g., é, ñ, ü)
- Non-Latin scripts (e.g., Cyrillic, Arabic, Chinese)
- Wastes the 8th bit in modern systems
- Inflexible for globalization, multilingual content, or emoji

Bits		b ₇	b ₆	b ₅	b ₄	b ₃	b ₂	b ₁	Column Row	0	0	0	1	0	1	0	0	1	1	0	0	1	0	1	1	0	1	1					
0	0	0	0	0	0	0	0	0	NUL	DLE	SP	0	@	P	`	p	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7	
0	0	0	0	1	1	1	1	1	SOH	DC1	!	1	A	Q	a	q	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7	
0	0	1	0	0	0	0	0	0	STX	DC2	"	2	B	R	b	r	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7	
0	0	1	1	0	0	0	0	0	ETX	DC3	#	3	C	S	c	s	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7	
0	1	0	0	0	0	0	0	0	EOT	DC4	\$	4	D	T	d	t	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7	
0	1	0	1	0	0	0	0	0	ENQ	NAK	%	5	E	U	e	u	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7	
0	1	1	0	0	0	0	0	0	ACK	SYN	&	6	F	V	f	v	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7	
0	1	1	1	0	0	0	0	0	BEL	ETB	'	7	G	W	g	w	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7	
1	0	0	0	0	0	0	0	0	BS	CAN	(8	H	X	h	x	1	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0
1	0	0	0	1	0	0	0	0	HT	EM)	9	I	Y	i	y	1	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0
1	0	1	0	0	0	0	0	0	LF	SUB	*	:	J	Z	j	z	1	0	1	0	0	0	0	0	1	0	0	0	0	0	0	0	0
1	0	1	1	0	0	0	0	0	VT	ESC	+	;	K	[k	{	1	0	1	1	0	0	0	0	1	0	0	0	0	0	0	0	0
1	1	0	0	0	0	0	0	0	FF	FS	,	<	L	\	l	{	1	1	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0
1	1	0	1	0	0	0	0	0	CR	GS	-	=	M]	m	}	1	1	0	1	0	0	0	0	1	0	0	0	0	0	0	0	0
1	1	1	0	0	0	0	0	0	SO	RS	.	>	N	^	n	~	1	1	1	0	0	0	0	0	1	0	0	0	0	0	0	0	0
1	1	1	1	0	0	0	0	0	SI	US	/	?	0	_	o	DEL	1	1	1	1	0	0	0	0	1	0	0	0	0	0	0	0	0

UTF-8 Encoding

- ❑ UTF-8: Unicode Transformation Format (8-bit)
- ❑ Encodes all Unicode characters using 1 to 4 bytes
- ❑ Compatible with ASCII — first 128 characters are identical
- ❑ Efficient for English, flexible for global languages
- ❑ Variable-length encoding:
 - 1 byte for basic Latin (e.g., A–Z)
 - 2–4 bytes for accented letters, symbols, and non-Latin scripts
- ❑ Self-synchronizing — easy to detect character boundaries
- ❑ Most common encoding on the web and in modern software

UTF-8 Examples

Character	Description	Unicode Code Point	UTF-8 Byte	UTF-8 Encoding (Hex)
A	Latin capital letter	U+0041	1 byte	41
é	Latin small letter e acute	U+00E9	2 bytes	C3 A9
你	Chinese character “you”	U+4F60	3 bytes	E4 BD A0
😊	Smiling face with smiling eyes	U+1F60A	4 bytes	F0 9F 98 8A

Byte Pair Encoding

- ❑ Most common method used today for sub-word encoding
- ❑ Encoder is trained on a large (often multi-lingual) corpus
- ❑ Start with a token vocabulary of single bytes [0-255] or 256 tokens
- ❑ Look for most common pair and add to the vocabulary
- ❑ Keep repeating until a desired maximum vocabulary
- ❑ Frequent sequences are encoded together
- ❑ GPT-2: ~50000 tokens developed in Feb 2019 by Open AI

BPE Encoding 1: Create the Dictionary

❑ Original text: ['a', 'a', 'a', 'b', 'd', 'a', 'a', 'a', 'b', 'a', 'c'] , Token dictionary: ['a', 'b', 'c', 'd']

❑ Step 1: Most common pair 'a'+a' (4 times)

- Merged to token →'aa'. New dictionary ['a', 'b', 'c', 'd', 'aa']
- Merged ['aa', 'a', 'b', 'd', 'aa', 'a', 'b', 'a', 'c']

❑ Step 2 Most common pair 'aa'+a' (or 'a'+b')

- Merge 'aa'+a' to 'aaa'. New dictionary: ['a', 'b', 'c', 'd', 'aa', 'aaa']
- New text: ['aaa', 'b', 'd', 'aaa', 'b', 'a', 'c']

❑ Step 3: Most common pair: ('aaa'+b')

- Merge to 'aaab'. New dictionary: ['a', 'b', 'c', 'd', 'aa', 'aaa', 'aaab']
- New text: ['aaab', 'd', 'aaab', 'a', 'c'].

❑ Step 4: Most common pair 'a', 'c' (there are others too)

- New dictionary: ['a', 'b', 'c', 'd', 'aa', 'aaa', 'aaab', 'ac']

BPE Encoding 2: Index the tokens

❑ Token dictionary from previous example:

- [‘a’, ‘b’, ‘c’, ‘d’, ‘aa’, ‘aaa’, ‘aaab’, ‘ac’]

❑ Assign each token an ID

Token	Token ID
a	0
b	1
c	2
d	3
aa	4
aaa	5
aaab	6
ac	7

BPE Encoding 3: Encode a New Sentence

- ❑ New example sentence: ‘aabacaadaaaabac’
- ❑ Find longest tokens possible
 - [‘aa’, ‘b’ ‘ac’, ‘aa’, ‘d’, ‘aaab’, ‘ac’]
- ❑ Set IDs
 - [4, 1, 7, 4, 3, 6, 7]
 -

Token	Token ID
a	0
b	1
c	2
d	3
aa	4
aaa	5
aaab	6
ac	7

Using a Pre-Trained Tokenizer

- ❑ Can download many tokenizers from Hugging Face

```
from transformers import AutoTokenizer

tokenizer = AutoTokenizer.from_pretrained("gpt2")

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://hug
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models o
    warnings.warn(
tokenizer_config.json: 100% [██████████] 26.0/26.0 [00:00<00:00, 2.04kB/s]
config.json: 100% [██████████] 665/665 [00:00<00:00, 58.2kB/s]
vocab.json: 100% [██████████] 1.04M/1.04M [00:00<00:00, 14.1MB/s]
merges.txt: 100% [██████████] 456k/456k [00:00<00:00, 19.6MB/s]
tokenizer.json: 100% [██████████] 1.36M/1.36M [00:00<00:00, 30.7MB/s]
```

Using the Tokenizer is Easy

```
example_text = "Hello world!"  
encoded_text = tokenizer.encode(example_text) [15496, 995, 0]  
  
print(encoded_text)
```

```
example_text2 = "Hello 😊世界!"  
encoded_text2 = tokenizer.encode(example_text2)  
  
print(encoded_text2) [15496, 30325, 232, 10310, 244, 45911, 234, 0]
```

Spaces are Part of the Token

- ❑ They are not removed
 - ❑ They are kept along with all other punctuation
-

```
text = ['world', ' world', 'World']
for t in text:
    enc = tokenizer.encode(t)
    token_id = enc[0]
    token = tokenizer.convert_ids_to_tokens(token_id)
    print('text: \'%s\' ID: %s token: %s' % (t, enc, token))
```

```
text: 'world' ID: [6894] token: world
text: ' world' ID: [995] token: Gworld
text: 'World' ID: [10603] token: World
```

Special character to represent
‘Space’ + ‘world’ is one token

Outline

❑ Tokenization

➡ Modern LLMs at a Glance

❑ Using a pre-trained LLM

Transformers

- ❑ Landmark paper, 2017
- ❑ Introduced the **transformer** architecture
- ❑ Originally for problems in NLP
- ❑ Now used in many fields:
 - ViT: An Image is Worth 16x16 Words,
Dosovitskiy et al, 2020
 - DETR End-to-End Object Detection with Transformers
Carion et al, 2020

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Lukasz Kaiser*
Google Brain
lukaszkaiser@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com

Prior Work

❑ Pre-2014 - RNN, LSTM, GRU



- RNNs struggled with long sequence data
- Could not handle long sequence lengths (forgetful)
- Example on next slide

❑ Neural Machine Translation by Jointly Learning to Align and Translate

- Bahdanau, 2014
- Landmark Paper

❑ Proposed New Architecture:

- Added attention mechanism to RNN
- Proposed initial idea of attention
- Model gave attention to particular hidden states when decoding



Challenges with Long-Term Dependencies

- ❑ NLP Challenge: **Long-term dependencies**
 - Language refers to items far in past
- ❑ Conventional architectures: RNN, LSTM, ...
 - Iteratively update state on each word / token
 - “Forget” about items in past
 - Struggled with language
- ❑ Need to consider a long-term **context**

Example Problem

When the **scientists** submitted their paper to the journal, they were hopeful. The **reviewers**, however, found several flaws in the methodology. After weeks of revisions and back-and-forth communication, the paper was finally accepted. But when the journal published it, **they** noticed that the **acknowledgments section** was missing, which caused quite a stir among the **authors**.

Question:

Who does “**they**” refer to?

Attention


❑ Takes two **contexts** of tokens

- **Query** and **key**
- Each $L \approx 1024$ length

$$A = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$


❑ Finds similarities $L \times L$ matrix

❑ Similarities look up L **values**




Attention Illustrated

- ❑ Find similarity matrix between tokens
- ❑ Example to the right:
 - Translation from French to English
- ❑ Can find long-term relations easily



Full Architecture


- ❑ Each layer consists of:
 - Multi-headed attention
 - Add and normalization
 - Feedforward neural network
- ❑ Repeated multiple attention layers
 - Skip connections between layers



Model Sizes

- ❑ Model are massive
- ❑ But more efficient architectures:
 - Retrieval-augmented-Generation (RAG)
 - Multi-Agent systems
 - Tool use and external APIs

Model	Estimated Parameters
GPT-4	~1.76 trillion
Claude 3	~2 trillion
Opus	
Gemini 1.5	~60–100 billion



Variant	Layers (n_layer)	Embedding Dim (n_embd)	Context Size (n_ctx)	Parameters
GPT-2 (small)	12	768	1024	117M
GPT-2 Medium	24	1024	1024	345M
GPT-2 Large	36	1280	1024	762M
GPT-2 XL	48	1600	1024	1.5B

Unsupervised Multitask Training

- ❑ Landmark OpenAI Paper
- ❑ Train on diverse tasks
 - Change final layers for each task
- ❑ Training requires:
 - 10000+ GPUs
 - Weeks + \$\$\$


The Llama 3 Herd of Models

Llama Team, AI @ Meta¹

GPUs	TP	CP	PP	DP	Seq. Len.	Batch size/DP	Tokens/Batch	TFLOPs/GPU	BF16 MFU
8,192	8	1	16	64	8,192	32	16M	430	43%
16,384	8	1	16	128	8,192	16	16M	400	41%
16,384	8	16	16	4	131,072	16	16M	380	38%

Language Models are Unsupervised Multitask Learners

Alec Radford *¹ Jeffrey Wu *¹ Rewon Child¹ David Luan¹ Dario Amodei **¹ Ilya Sutskever **¹



Pre-Trained Models in Hugging Face

❑ Hugging Face

- Has 1000s of pre-trained models
- Sentence Transformers


❑ Easy-to-use

- No training
- Varying complexity

The screenshot shows the Hugging Face library interface. At the top, there's a search bar labeled "Filter Libraries by name" and a "Reset Libraries" button. Below the search bar is a row of library logos: PyTorch, TensorFlow, JAX, Safetensors, Transformers, PEFT, TensorBoard, GGUFG, Diffusers, ONNX, and stable-baselines3. Further down are rows of ML-related libraries: sentence-transformers, ml-agents, MLX, MLX, TF-Keras, Keras, Adapters, setfit, timm, Transformers.js, sample-factory, Joblib, and OpenVINO. To the right of the interface, a grid of 12 pre-trained model cards is displayed:

- google/embeddinggemma-300m: Sentence Similarity, 0.3B, Updated Sep 25, 428k, 1.21k
- sentence-transformers/all-MiniLM-L6-v2: Sentence Similarity, 22.7M, Updated Mar 6, 144M, 4.12k
- Owen/Owen3-Embedding-0.6B: Feature Extraction, 0.6B, Updated Jun 20, 4.96M, 738
- nvidia/llama-embed-nemotron-8b: Feature Extraction, 8B, Updated 3 days ago, 102k, 73
- tencent/KaLM-Embedding-Gemma3-12B-2511: Sentence Similarity, 12B, Updated 13 days ago, 8.63k, 19
- infgrad/Jasper-Token-Compression-600M: Sentence Similarity, 0.6B, Updated 1 day ago, 61, 9
- sentence-transformers/all-mppnet-base-v2: Sentence Similarity, 0.1B, Updated Aug 19, 19M, 1.19k
- Alibaba-NLP/gme-Owen2-VL-2B-Instruct: Sentence Similarity, 2B, Updated Jun 9, 87.3k, 118
- intfloat/multilingual-e5-large: Feature Extraction, 0.6B, Updated Feb 16, 2.64M, 1.08k
- BAAI/bge-large-en-v1.5: Feature Extraction, 0.3B, Updated Feb 20, 2024, 5.24M, 59

Outline

- ❑ Tokenization
- ❑ Modern LLMs at a Glance
-  Using a pre-trained LLM

Movie Sentiment Analysis

- ❑ Problem: Given text for a movie detect sentiment?
 - Not given the numeric score
- ❑ Requires deep understanding of text

< The Lord of the Rings: The Rings of... ...

★ 2
Bland and mediocre
I would have loved this to be a really good sword and sorcery series. Its barely even average. And when you think how much money Amazon have spent on...



“Negative review”

IMDB Review Data

- ❑ Dataset from 2011 classic paper
- ❑ 50,000 reviews taken from IMBD database
- ❑ Available from Kaggle

Learning Word Vectors for Sentiment Analysis

**Andrew L. Maas, Raymond E. Daly, Peter T. Pham, Dan Huang,
Andrew Y. Ng, and Christopher Potts**

Stanford University
Stanford, CA 94305

[amaas, rdaly, ptpham, yuze, ang, cgpotts]@stanford.edu

```
import kagglehub

# Download latest version
path = kagglehub.dataset_download("lakshmi25npathi/imdb-dataset-of-50k-movie-reviews")
```

Example Reviews

As another reviewer states Hanna's War is an outstanding film about an outstanding person, Hanna "Anniko" Senesh, who would become the Jewish Joan Of Arc. Unfortunately I diverge in opinion not agreeing that Miss Detmers as the lead is too beautiful to be taken seriously as a resistance fighter. In truth for me her performance is not held back by her beauty but makes it all the more stark in the terror of the sadistic brutality as a resistor she faces. Maruschka Detmers performance is brave, poignant, heartfelt or understood, and totally believable. In other words for me "In the zone." from the opening credits. If you would like to learn about the suffering of someone else for something they believe in and be impressively entertained give Hanna's War with Maruschka Detmers a try. My hat is off also to Ellen Burstyn as Hanna's mother a much well known and famous actress who could have made effort to walk off with the film. In that it is a team effort perhaps of two actress' but not an All About Eve situation.

Example Positive Review

This film is so bad - dialogues, story, actors and actresses - everything! - that it's hard to imagine that we'll see a worse movie this year or in the following years. "Love's Brother" (set in Australia among Italian immigrants) has nothing but shallow clichés about Italian culture to offer, and it is quite telling that even the Italians from and in Italy speak ENGLISH in the film. The message of the film - ugly people have to marry ugly people, beautiful people have to marry beautiful people - is truly discomforting. Giovanni Ribisi is quite good in films like 'Suburbia' or 'Lost in Translation', but here his pseudo-Italian accent is hard to bear. See this film at your own risk. Trash as trash can!

Example Negative Review

Downloading a Pre-Trained LLM

❑ Use MiniLMV2:

- From Microsoft 2021
- Very compact, easy to use

❑ Download from Hugging Face

**MINILMv2: Multi-Head Self-Attention Relation Distillation
for Compressing Pretrained Transformers**

Wenhui Wang Hangbo Bao Shaohan Huang Li Dong Furu Wei*
Microsoft Research
{wenwan,t-habao,shaohanh,lidong1,fuwei}@microsoft.com

```
from sentence_transformers import SentenceTransformer


model = SentenceTransformer('sentence-transformers/all-MiniLM-L6-v2')
model = model.to(device)
```

Encoding the Reviews

- ❑ Encode the reviews with the model
- ❑ On Google Colab free T4 GPU
 - About 2 mins for 50,000 reviews
 - With CPU it is too long
- ❑ Each embedding is dimension 384
- ❑ Output a 50000×384 matrix



50000 reviews



```
review_embeddings = model.encode(reviews, convert_to_tensor=True, show_progress_bar=True)
print(f"Shape of review embeddings: {review_embeddings.shape}")
print(f"First review embedding (partial):\n{review_embeddings[0][:5]}")
```

Training a Classifier

- Used a simple neural network
 - Input = 384 (embedding dimension)
 - 50 hidden units
 - Binary output (positive or negative)

```
class SentimentClassifier(nn.Module):  
    def __init__(self, input_size, hidden_size, num_classes):  
        super(SentimentClassifier, self).__init__()  
        self.fc1 = nn.Linear(input_size, hidden_size)  
        self.relu = nn.ReLU()  
        self.fc2 = nn.Linear(hidden_size, num_classes)  
        self.sigmoid = nn.Sigmoid()  
  
    def forward(self, x):  
        out = self.fc1(x)  
        out = self.relu(out)  
        out = self.fc2(out)  
        out = self.sigmoid(out)  
        return out  
  
# Determine input size from the embeddings  
input_size = review_embeddings.shape[1]  
hidden_size = 200  
num_classes = 1 # For binary classification
```

Classification Results

❑ Not the best, but it works

